

SI2021

MNIST and Tensorboard hands on
Paul Rodriguez, PhD
SDSC



Overview: MNIST w/noise, and tensorboard

MNIST is too easy, so I switched labels to make learning harder

Tasks:

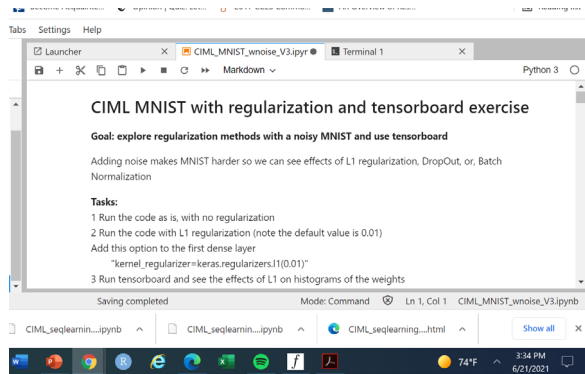
- 1 Try adding L1 regularization in classification layer (a penalty $\lambda * \sum |W|$ is added to Loss)
- 2 View results in tensorboard

Optional, try adding these and comparing performance in tensorboard:

dropout layer (drop % of activations of a layer)

batchNormalization (center and scale activations of a layer)

1 Read the instructions and glance at the code sections



Notice there is a “logs” directory set up for tensorflow

Notice there is a new callback function to log information

Next cell sets up directory for tensorboard logs

```
[1]: #Set up the location for tensorflow logs
import datetime, os
logdir = os.path.join("logs",
    datetime.datetime.now().strftime("%Y%m%d-%H%M%S")+'_' + \
    'test') #<<<<----- you can add a comment to describe the test
print('using',logdir,' for logs')
```

```
5]: from tensorflow.keras.callbacks import EarlyStopping
ES_function = EarlyStopping(monitor='val_loss', mode='min', patience=20)
    #or try: monitor='val_accuracy', mode='max', patience=1)

tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1)
    #Use this to get performance info, like memory usage: profile_batch = '5,20')

fit_history=mymodel.fit(X_train, Y_train, #validation_split=0.20,
    validation_data=(X_test,Y_test),
    batch_size=64, epochs=100, verbose=1,callbacks=[ES_function,tensorboard_callback])
```

2 Run notebook , following instructions

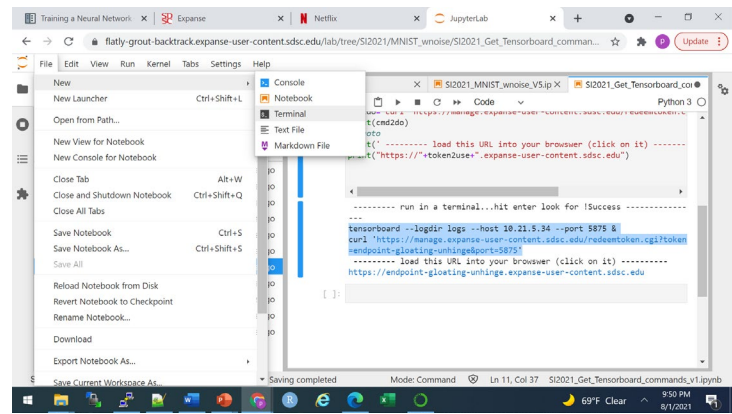
On Expanse, we set up tensorboard to use reverse proxy as follows:

1. Open a terminal window, cd to SI2021 4.1b DL session
2. Open SI2021_Get_Tensorboard_commands notebook and run it

It should print out a 'tensorboard' command and 'curl' command, and a URL

3. cut & paste commands into terminal window

4. Click on URL



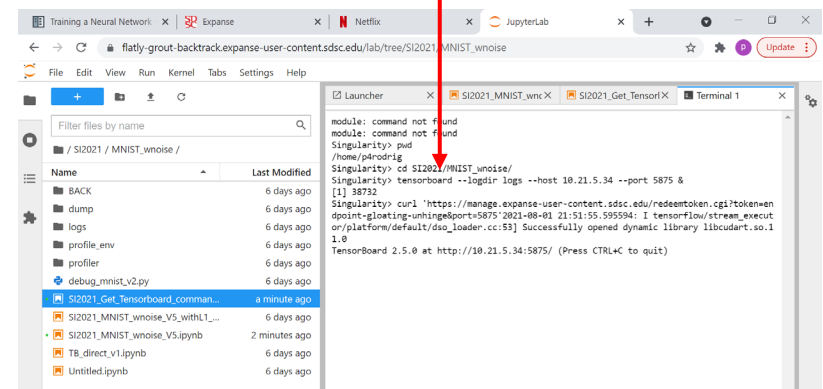
SI2021_Get_Tensorboard_commands_v1 (autosaved)

```
import os
ip2use = getnet hosts $(hostname -s).cm.cluster | awk '{print $1}'
token2use_res = curl https://manage.expense-user-content.sdsc.edu/getlink.cgi?token2use=
token2use = token2use_res[-1]
port2use = str(int(np.random.uniform(4000,6000,1)))

#2 Output command for user to enter into terminal window
print('----- run in a terminal...hit enter look for !Success -----')
print('tensorboard --logdir logs --host ' + ip2use[0] + ' --port ' + port2use + ' &')
#3 redeem token with port
cmd2do = curl 'https://manage.expense-user-content.sdsc.edu/redeemtoken.cgi?token=' + token2use + '&port=' + port2use + ''
print(cmd2do)
#4 goto
print('----- load this URL into your browser (click on it) -----')
print('https://' + token2use + '.expense-user-content.sdsc.edu')
```

----- run in a terminal...hit enter look for !Success -----
tensorboard --logdir logs --host 10.21.3.31 --port 4125 &
curl 'https://manage.expense-user-content.sdsc.edu/redeemtoken.cgi?token=oaat-pelican-regain&port=4125'
----- load this URL into your browser (click on it) -----
https://oaat-pelican-regain.expense-user-content.sdsc.edu

cut & paste

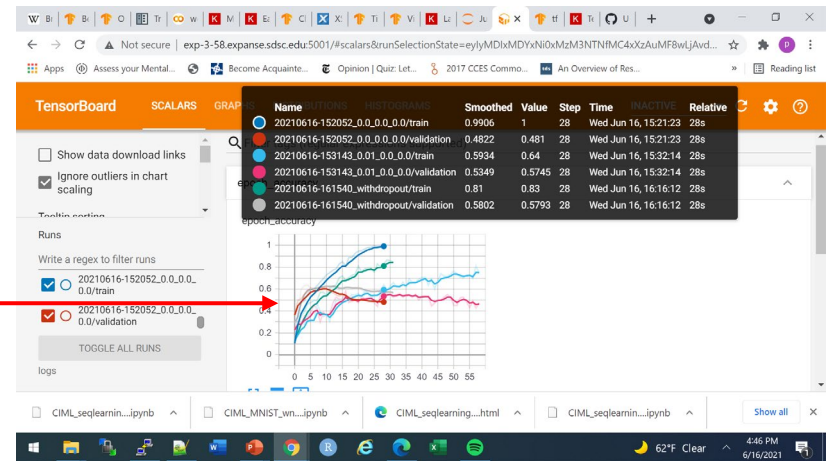
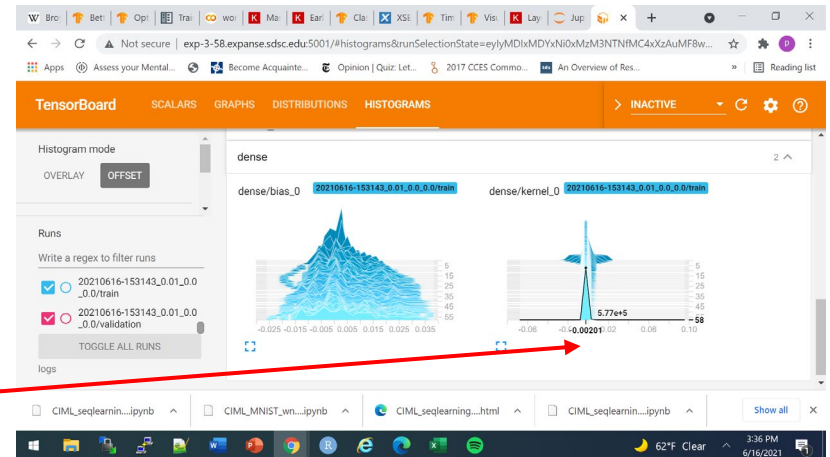


Run the notebook;

1 Find the <<<< ---- comments to change/add regularization

2 In tensorboard, toggle the runs to just see the logs with L1 regularization, then check out the weight histograms dense layer

3
If you run several times you can get all the performances plotted together (depending on what log data you toggle)



NOTE on optional tasks:


BatchNormalization seems
to learn the training set
very fast and avoid the
inverted U for accuracy

```
numfilters = 64
mymodel.add(Convolution2D(numfilters, (3,3), strides=1, data_format="channels_last", activation='relu')
mymodel.add(Convolution2D(numfilters, (3,3), strides=1, data_format="channels_last", activation='relu')
mymodel.add(MaxPooling2D(pool_size=(2,2),strides=2,data_format="channels_last"))
mymodel.add(Flatten())

#----- add final classification layers
mymodel.add(Dense(64, activation='relu')) #<<<<----- Add the L1 reglzer option here
#mymodel.add(Dropout(0.50))
mymodel.add(BatchNormalization(axis=-1))

mymodel.add(Dense(10, activation='softmax'))

print('added layers to model')
```



PAUSE

EXTRA NOTE:

Tensorboard has a 'debugger' option

The debugger will 'alert' you to errors, such as NaNs in your calculations. Most useful if you are building custom functions.

It requires a command in your Python code to 'dump' information into the logs

```
tf.debugging.experimental.enable_dump_debug_info( ....
```



EXTRA NOTE:

Tensorboard has a 'profiler' plugin

The profile option has performance information for some part of the training iterations

On Expanse it requires installation of virtual environment (ask for details if interested)

